

DR. FRANCISCO PINA-MARTINS (Orcid ID : 0000-0003-1836-397X)

DR. DIOGO NUNO SILVA (Orcid ID : 0000-0002-6827-2673)

Article type : Resource Article

**Title:** *Structure\_threader*: an Improved Method for Automation Parallelization of Programs STRUCTURE, FASTSTRUCTURE and *Maverick* on Multi Core CPU systems.

**Authors:** Pina-Martins, Francisco<sup>ab</sup>; Silva, Diogo N<sup>ad</sup>; Fino, Joana<sup>ac</sup>; Paulo, Octávio S.<sup>a</sup>

**Addresses:**

<sup>a</sup>Computational Biology and Population Genomics Group, Centre for Ecology, Evolution and Environmental Changes, Departamento de Biologia Animal, Faculdade de Ciências, Universidade de Lisboa, Campo Grande, 1749-016 Lisboa, Portugal

<sup>b</sup>Departamento de Biologia e CESAM, Univ. de Aveiro, Portugal

<sup>c</sup>Department of Human Genetics, National Institute of Health Dr Ricardo Jorge, Lisbon, Portugal

<sup>d</sup>Instituto Superior de Agronomia, Universidade de Lisboa, CIFIC - Centro de Investigação das Ferrugens do Cafeeiro, Oeiras, PT.

**Keywords:** Bioinformatics/Phyloinformatics, Genomics/Proteomics, Molecular Evolution, Population Genetics – Empirical, Clustering, Parallel computing

**Corresponding author:** Francisco Pina-Martins; Faculdade de Ciências da Universidade de Lisboa, Campo Grande, Ed. C2, Sala 2.3.22, 1749-016 Lisboa, Portugal; f.pinamartins@gmail.com

**Running title:** *Structure\_threader*

## 1. Abstract

*Structure\_threader* is a program to parallelize multiple runs of genetic clustering software that does not make use of multi-threading technology (STRUCTURE, FASTSTRUCTURE and *Maverick*) on multi-core computers. Our approach was benchmarked across multiple systems and displayed great speed improvements relative to the single threaded implementation, scaling very close to linearly with the number of physical cores used.

This article has been accepted for publication and undergone full peer review but has not been through the copyediting, typesetting, pagination and proofreading process, which may lead to differences between this version and the Version of Record. Please cite this article as doi: 10.1111/1755-0998.12702

This article is protected by copyright. All rights reserved.

*Structure\_threader* was compared to previous software written for the same task - *ParallelStructure* and *StrAuto*, and was proven to be the faster (up to 25% faster) wrapper under all tested scenarios.

Furthermore, *Structure\_threader* can perform several automatic and convenient operations, assisting the user in assessing the most biologically likely value of 'K' via implementations such as the "Evanno", or "Thermodynamic Integration" tests and automatically draw the "meanQ" plots (static or interactive) for each value of K (or even combined plots).

*Structure\_threader* is written in python 3 and licensed under the GPLv3. It can be downloaded free of charge at [https://github.com/StuntsPT/Structure\\_threader](https://github.com/StuntsPT/Structure_threader).

## 2. Introduction

Clustering analyses are widely used in population genetics and, nowadays, population genomics. This technique of using multilocus genotype data to infer population clusters, is frequently performed based on multiple MCMC re-sampling. One of the most popular tools for performing this type of analyses is STRUCTURE (Pritchard, Stephens, & Donnelly, 2000). Despite producing robust results, this approach demands long run times, even in modern machines. This problem is aggravated as the type of analysed datasets, which gradually grow from relatively small, such as microsatellite loci (De Barro, 2005; Muchadeyi et al., 2007), to high throughput sequencing (Lamaze, Sauvage, Marie, Garant, & Bernatchez, 2012; Renaut, Grassa, Moyers, Kane, & Rieseberg, 2012), consequently increasing run times by orders of magnitude.

The process can be sped up by either running multiple instances of the used software, which is an inefficient and error prone method requiring constant attention and intervention from the user. There are faster software alternatives to STRUCTURE, which can also be used to speed up the analysis process.

One such option is analysing the data in the program FASTSTRUCTURE (Raj, Stephens, & Pritchard, 2014), which decreases run times by up to two orders of magnitude. However FASTSTRUCTURE does not support the popular "no admixture" model present in STRUCTURE, and is not capable of handling

haploid data (and several other less widely used features), which limits its application to a wide range of data.

Another option in to analyse the data in the software *Maverick* (Verity & Nichols, 2016), which is also considerably faster than STRUCTURE, but not as fast as FASTSTRUCTURE by an order of magnitude. It does, however support most of the same features as STRUCTURE and uses a built-in, improved method for helping determine the most biologically relevant value of “K” called “Thermodynamic Integration” (Verity & Nichols, 2016). Regardless of the speed gains these programs offer, they are only able to use a single CPU core for their computations, which means that these methods too, do not scale well with current multi-core IT infrastructure.

Alternatively, a method to bootstrap multiple simultaneous runs of the software can be used, such as the R (R Core Team, 2013) package *ParallelStructure* (Besnier & Glover, 2013), or *StrAuto* (Chhatre & Emerson, 2017), which does exactly that for the software STRUCTURE (Pritchard et al., 2000). *ParallelStructure*, however, has scaling problems, as described in the manuscript, considerably losing efficiency as more CPU cores are used. *StrAuto* is another option that does indeed scale well with the number of CPU cores used, but like *ParallelStructure*, it only works as a wrapper for the software STRUCTURE, and cannot be used to speed up other popular genetic clustering programs.

Furthermore, after the clustering step is finished, it is necessary to infer the number of clusters that make most biological sense for the data (Earl & vonHoldt, 2012), using methods such as the “Evanno test” (Evanno, Regnaut, & Goudet, 2005), or the “Thermodynamic Integration” (TI) method (Verity & Nichols, 2016). After this, it is also often necessary to plot the “meanQ” values of each cluster per individual, to be able to interpret the biological significance of the data. This is usually done with software such as DISTRUCT (ROSENBERG, 2004).

All of these steps typically require parsing the results files of each clustering run and manually running all the required steps until the final outcome is produced (Earl & vonHoldt, 2012). This is not only time consuming as it is also error prone due to the large number of separate steps that must be

taken during the process. Neither *ParallelStructure* nor *StrAuto* provide an automated and reproducible way to perform this task.

Part of this process is largely facilitated by the program STRUCTURE HARVESTER (Earl & vonHoldt, 2012), which automates the parsing of STRUCTURE runs and uses that information to perform an “Evanno test” on the data, which uses some heuristics to predict which value of ‘K’ makes the most biological sense regarding the analysed data. Although this is a very convenient automation, it still relies on manual user intervention to input the data from STRUCTURE, does not provide assistance with the plotting of the “meanQ” values and only works for the software STRUCTURE. Other programs, such as FASTSTRUCTURE include the necessary software to perform these tests, and even to plot the “meanQ” values, but still require manual intervention between these steps. *Maverick* goes further and presents the full posterior distribution for ‘K’ using the “Thermodynamic Integration” test as an automatic last step of the analysis and even recommends some scripts for drawing “meanQ” plots, but this last step also requires human intervention.

To address these two problems (reducing run times and automating the analyses tasks), we herein present *Structure\_threader*: a program to parallelize STRUCTURE, FASTSTRUCTURE and *Maverick* runs that considerably reduces the scaling problems of previous approaches and automates the entire process, - wrapping the runs, assisting in the choice of the most biologically relevant value of K, and drawing the “meanQ” plots.

*Structure\_threader* is available on [https://github.com/StuntpT/Structure\\_threader](https://github.com/StuntpT/Structure_threader). For the stable (non development) versions, check the releases page, or get it from Pypi.

### 3. Materials & Methods

#### 3.1. Program description

*Structure\_threader*, licensed under the GPLv3, is an open source program written in python (<https://www.python.org/>) that automates and parallelizes genetic clustering software (STRUCTURE, FASTSTRUCTURE and *Maverick*) runs.

The software was written according the “Best Practices in the Development of Bioinformatics Software” (Leprevost, Barbosa, Francisco, Perez-Riverol, & Carvalho, 2014) and can be run on any platform where python is available, such as GNU/Linux, Mac OS X and Microsoft Windows (other platforms may also work, but were not tested). Additional details are available in the program's documentation.

All options supported by the wrapped programs can be passed to *Structure\_threader* as command line arguments. These are explained in detail in both the program's online documentation and builtin help text.

Parameters are passed to the wrapped software as in their default implementations – all wrapped programs take arguments from the command line, STRUCTURE also reads settings from the files “mainparams” and “extraparams” and *Maverick* from “parameters.txt”.

After performing the parallelized runs of the wrapped software, *Structure\_threader* runs a slightly modified (for integrating with *Structure\_threader*) version of STRUCTURE HARVESTER, *chooseK.py* (Raj et al., 2014) or TI for helping identify the most biologically relevant value of ‘K’ for any given dataset.

Finally, *Structure\_threader* parses the results files and draws the “meanQ” plots for each considered value of “K”. These are drawn in both an interactive version for visualization and in a static version, better suited for publication.

Example data files and results are provided in the program's repository.

### 3.2. Threading strategy

The threading strategy used in *Structure\_threader* is represented in Figure 1. *Structure\_threader* takes the provided input file, the values of “K” to test and the required number of replicates, and creates a job queue, which is sorted by decreasing complexity order. After this,  $P$  child processes are spawned, (where  $P$  is the number of threads made available to the software) each containing one independent instance of the wrapped program. Each of these child processes takes the first available job from the queue and once it is finished, its output is processed by the main process for error handling and logging. The child processes are spawned using python’s “multiprocessing” and “subprocess” modules from the standard library.

### 3.3. Benchmarking process

In order to assess the gains provided by *Structure\_threader*, the program was benchmarked in four different systems, described in Table 1, with various specifications. Runs were performed twice to serve as replicates (Supplementary Table 1). Run times for STRUCTURE were assessed using both *Structure\_threader* v0.4.1, *ParallelStructure* v1.0 and *StrAuto* v1.0, which were then compared. FASTSTRUCTURE and *Maverick* runs were only wrapped in *Structure\_threader*, since none of the other programs supports this, and compared with the default, single threaded implementation.

Usage of RAM was monitored during the benchmarking process, and it was never detected as a bottleneck on any of the systems. None of the wrapped programs is very I/O intensive (at least as far as the tested systems were concerned), meaning that the present tests were always CPU bound.

Run times were measured using *zsh*’s *time* builtin method (wall time was used), and then normalized to a “speed up” factor (Besnier & Glover, 2013) by dividing the time of the multi-core runs by the time of the single core runs, which were performed in the measured programs’ default implementations.

## Benchmarking STRUCTURE

### 3.3.1. Test dataset description

The test file used for the benchmarks consists of 100 individuals, represented by 80 SNP loci without missing data. This dataset was crafted based on data from the *1000 genomes project* (The 1000 Genomes Project Consortium, 2015) to perform the benchmarks and was constructed aiming for a size that would be neither too small, which could bias the benchmarking towards very quick runs, nor too large, to avoid the process taking too long to be practical.

This dataset was created from public data, and instructions on how to recreate it are available in the documentation and in the program's repository.

### 3.3.2. Benchmark details

The benchmarking process consisted of running the test dataset on STRUCTURE v2.3.4 for  $1 \times 10^6$  MCMC iterations, and a “burnin” length of  $5 \times 10^4$ , under “admixture” model (for other parameters check the program’s repository). These settings were performed for values of “K” varying from 1 to 4. Each value of “K” was run with 4 replicates, which means a total of 16 STRUCTURE runs were performed in each benchmark. All these runs were performed on the default, single threaded implementation and under the *Structure\_threader*, *ParallelStructure* (using the “parallel\_structure()” implementation, which initial testing found to be faster in all used machines) and *StrAuto* wrappers.

## 3.4. Benchmarking FASTSTRUCTURE

### 3.4.1. Test dataset description

The test dataset used for benchmarking FASTSTRUCTURE runs, is different from the one used for benchmarking STRUCTURE, since this program was designed to analyse larger datasets. The tested file consists of 1000 individuals, represented by 1000 SNP loci. Like the previous dataset, this one was

also crafted from the same public data from the 1000 genomes project, and instructions for recreating it are available in the documentation. The used dataset itself is available in the program's repository.

### 3.4.2. Benchmark details

The benchmarking process consisted of running the above described dataset for values of "K" from 1 to 16 for each benchmark run.

The average run time of both replicates was used to plot and analyse the data. Since a FASTSTRUCTURE runs do not require replicates for downstream analyses, each value of "K" was run only once per benchmark, which means that a total of 16 FASTSTRUCTURE runs were performed both in the default implementation and under the *Structure\_threader* wrapper.

## 3.5. Benchmarking *Maverick*

### 3.5.1. Test dataset description

The test file used for the *Maverick* benchmarks is the same that was used to benchmark STRUCTURE, which is described above.

### 3.5.2. Benchmark details

The benchmarking process consisted of running the test dataset on *Maverick* v1.0.4 for  $1 \times 10^4$  MCMC iterations, plus a "burnin" length of 2500 iterations, with 5 replicates each (for other parameters check the program's repository). These settings were performed for values of "K" varying from 1 to 16.



## 4. Results & Discussion

Using *Structure\_threader* as a wrapper for all wrapped programs has yielded increases in speed that scales almost linearly with the number of processes used, at least as long as physical cores are concerned, as can be seen in Figure 2 and Figure 3.

Considering the benchmark results present in Figure 2, it is clear that both *Structure\_threader* and *StrAuto* are more efficient methods to perform multiple STRUCTURE runs on multiple core systems than *ParallelStructure* (on average 7% faster in the tested systems, varying from 1% to 25% faster). *Structure\_threader* is also always faster than *StrAuto*, but by much smaller margins than when compared with *ParallelStructure* (on average 3% faster, varying from 0.3% to 7% faster). Regardless of the tested system and number of cores used, the differences in “speed up” are always in favour of *Structure\_threader*. When compared to *ParallelStructure*, the difference increases with the requested scaling – the more physical cores are used, the better the relative performance of *Structure\_threader*. Also worth noting is that the “speed up” values obtained here with *ParallelStructure* when using physical cores, are somewhat better than what is described in (Besnier & Glover, 2013), but this could be due to differences in benchmark workloads.

Speed up differences between *StrAuto* and *Structure\_threader* are small, but can be compared in Figure 2. A more detailed comparison, can be made using the data tables available in Supplementary Material 1.

Unlike *ParallelStructure* and *StrAuto*, *Structure\_threader* can also speed up the runs of other programs by running them in parallel. Similar to what is done for running STRUCTURE, wrapping FASTSTRUCTURE and *Maverick* in *Structure\_threader*, provides considerable speed improvements, once again scaling almost linearly as long as *hyper-threading* is not in effect (Figure 3).

Although ideally the “speed up” factor should scale linearly with the number of used physical cores, this does not always happen in practice (Figure 2 and Figure 3). Of the three tested wrappers, *Structure\_threader* scales the closest to linearly, even when using 8 physical cores, where the “speed

up” factor varies between 6.24 and 7.91, depending on the system and the wrapped program.

*ParallelStructure* shows the worst scaling of the tested wrappers, especially on 8 physical threads, where the “speed up” factor varies between 5.95 and 6.85.

The large drop in performance increase, regardless of the used wrapper program, happens when using *hyper-threading* (using more than eight cores in the *Nehalem Rack* system and more than four in the *Haswell Desktop* system – the CPUs of the other two systems do not have this feature), as is sometimes described under certain workloads (Leng, Ali, Hsieh, Mashayekhi, & Rooholamini, 2002; Marr et al., 2002; Saini et al., 2011). We are not sure why this happens on this particular workload, but the issue is not as evident when analysing smaller datasets as the one from (Besnier & Glover, 2013). It is therefore hypothesised that it may be related to “cache thrashing”, a phenomenon that occurs when the CPU constantly refreshes the contents of L2 and L3 caches for quickly accessing different information. “Cache thrashing” is more likely to happen when working with larger datasets and when *hyper-threading* is active since both logical cores share L2 and L3 cache (Marr et al., 2002).

The automated plot drawing feature of *Structure\_threader* is responsible for both a simplification of the analysis process (less steps per analysis), and also for the reduction in random error (consequence of less human intervention).

The mentioned plots produced by *Structure\_threader* are provided in two formats. A static, vectorial image in “svg” format, especially suited for publication, and an interactive HTML version of the plot, suited for results exploration.

## 5. Conclusions

The observed difference in efficiency between *Structure\_threader* and *ParallelStructure* can probably be explained by the programming languages utilized in the wrappers (Python vs. R) and the fact that *ParallelStructure* solves tasks in increasing order of complexity, whereas *Structure\_threader* sorts them in decreasing order. This strategy provides an increase in efficiency, since the sorting minimizes the time each CPU core is left idle.

Another important difference between *ParallelStructure* and *Structure\_threader* is that the former is a framework to build scripts that perform the requested analyses, and the latter can either be used as a framework, or directly from the command line. This makes *Structure\_threader* much easier to use, while simultaneously keeping the same type of flexibility *ParallelStructure* offers. Although both programs can be used to draw the clustering plots from the STRUCTURE results, the features offered by *Structure\_threader* go far beyond the basic plotting that *ParallelStructure* is capable of.

The speed gains obtained with *Structure\_threader* and *StrAuto* are very similar, with only a marginal difference favouring *Structure\_threader*. This difference is likely due to the efficiency of *python's* higher speed when compared to *bash's*, and eventually due to a smaller overhead of *python's multiprocessing* module when compared to that of GNU parallel (Tange, 2011). Although both programs are run from the command line interface, *Structure\_threader* is more user friendly than *StrAuto*, since it includes built-in help, handles user errors, and allows for a lot of parameters to be defined directly in the command line.

*Structure\_threader* was designed to exploit the power of multi-core machines for speeding up multiple genetic clustering software runs, with emphasis on scalability. Our results demonstrate that in every tested scenario this goal is fulfilled in a more efficient way than previous solutions.

Furthermore, *Structure\_threader* goes much farther than the two previous solutions in its capabilities to perform tests for estimating the most biologically relevant “K” value, as well as plotting flexibility.

Although the automation process that *Structure\_threader* provides does not decrease computation time, it should significantly speed up the analyses process, due to the human time that is saved. Furthermore, this automation is also one important step for reproducibility of the studies that use this software. That being said, it is also important that users interact with and explore the options and parametrization the wrapped programs offer. It is critical that these are well understood in order to obtain meaningful and statistically relevant results.

We find that the obtained decrease in run times, allied with the ease of use and automation, including that of follow up analysis, make *Structure\_threader* a useful tool to any investigator working with population genetics/genomics data and the best current choice for performing genetic clustering analyses.

## 6. Acknowledgements

This study was financed by Portuguese National Funds, through FCT – Fundação para a Ciência e a Tecnologia, within the projects UID/BIA/00329/2013, SFRH/BD/51411/2011 and SFRH/BD/86736/2012.

We would further like to thank Bob Verity and an anonymous reviewer for their suggestions that thoroughly improved both the software and the manuscript.

## References

- Besnier, F., & Glover, K. A. (2013). ParallelStructure: A R Package to Distribute Parallel Runs of the Population Genetics Program STRUCTURE on Multi-Core Computers. *PLoS ONE*, 8(7).  
<https://doi.org/10.1371/journal.pone.0070651>
- Chhatre, V. E., & Emerson, K. J. (2017). StrAuto: automation and parallelization of STRUCTURE analysis. *BMC Bioinformatics*, 18, 192. <https://doi.org/10.1186/s12859-017-1593-0>
- De Barro, P. J. (2005). Genetic structure of the whitefly *Bemisia tabaci* in the Asia–Pacific region revealed using microsatellite markers. *Molecular Ecology*, 14(12), 3695–3718.  
<https://doi.org/10.1111/j.1365-294X.2005.02700.x>
- Earl, D. A., & vonHoldt, B. M. (2012). STRUCTURE HARVESTER: a website and program for visualizing STRUCTURE output and implementing the Evanno method. *Conservation Genetics Resources*, 4(2), 359–361. <https://doi.org/10.1007/s12686-011-9548-7>
- Evanno, G., Regnaut, S., & Goudet, J. (2005). Detecting the number of clusters of individuals using the software structure: a simulation study. *Molecular Ecology*, 14(8), 2611–2620.  
<https://doi.org/10.1111/j.1365-294X.2005.02553.x>
- Lamaze, F. C., Sauvage, C., Marie, A., Garant, D., & Bernatchez, L. (2012). Dynamics of introgressive

hybridization assessed by SNP population genomics of coding genes in stocked brook charr (*Salvelinus fontinalis*). *Molecular Ecology*, 21(12), 2877–2895.

<https://doi.org/10.1111/j.1365-294X.2012.05579.x>

Leng, T., Ali, R., Hsieh, J., Mashayekhi, V., & Rooholamini, R. (2002). An Empirical Study of Hyper-Threading in High Performance Computing Clusters. In *Linux Clusters: The HPC Revolution, 2002*. St. Petersburg.

Leprevost, F. da V., Barbosa, V. C., Francisco, E. L., Perez-Riverol, Y., & Carvalho, P. C. (2014). On best practices in the development of bioinformatics software. *Bioinformatics and Computational Biology*, 5, 199. <https://doi.org/10.3389/fgene.2014.00199>

Marr, D., Binns, F., Hill, D. L., Hinton, G., Koufty, D. A., Miller, A. J., & Upton, M. (2002). Hyper-Threading Technology Architecture and Microarchitecture. *Intel Technology Journal*, 6. Retrieved from [https://www.researchgate.net/publication/237005389\\_Hyper-Threading\\_Technology\\_Architecture\\_and\\_Microarchitecture](https://www.researchgate.net/publication/237005389_Hyper-Threading_Technology_Architecture_and_Microarchitecture)

Muchadeyi, F. C., Eding, H., Wollny, C. B. A., Groeneveld, E., Makuza, S. M., Shamseldin, R., ... Weigend, S. (2007). Absence of population substructuring in Zimbabwe chicken ecotypes inferred using microsatellite analysis. *Animal Genetics*, 38(4), 332–339. <https://doi.org/10.1111/j.1365-2052.2007.01606.x>

Pritchard, J. K., Stephens, M., & Donnelly, P. (2000). Inference of population structure using multilocus genotype data. *Genetics*, 155(2), 945–959.

R Core Team. (2013). *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from <http://www.R-project.org/>

Raj, A., Stephens, M., & Pritchard, J. K. (2014). fastSTRUCTURE: Variational Inference of Population Structure in Large SNP Data Sets. *Genetics*, 197(2), 573–589. <https://doi.org/10.1534/genetics.114.164350>

Renaut, S., Grassa, C. J., Moyers, B. T., Kane, N. C., & Rieseberg, L. H. (2012). The Population Genomics of Sunflowers and Genomic Determinants of Protein Evolution Revealed by RNAseq. *Biology*, 1(3), 575–596. <https://doi.org/10.3390/biology1030575>

Rosenberg, N. A. (2004). distruct: a program for the graphical display of population structure. *Molecular Ecology Notes*, 4(1), 137–138. <https://doi.org/10.1046/j.1471-8286.2003.00566.x>

Saini, S., Jin, H., Hood, R., Barker, D., Mehrotra, P., & Biswas, R. (2011). The impact of hyper-threading on processor resource utilization in production applications. In *2011 18th International Conference on High Performance Computing* (pp. 1–10).

<https://doi.org/10.1109/HiPC.2011.6152743>

Tange, O. (2011). GNU Parallel—The Command-Line Power Tool. *Login: The USENIX Magazine*, 36(1), 42–47.

The 1000 Genomes Project Consortium. (2015). A global reference for human genetic variation. *Nature*, 526(7571), 68–74. <https://doi.org/10.1038/nature15393>

Verity, R., & Nichols, R. A. (2016). Estimating the Number of Subpopulations (K) in Structured Populations. *Genetics*, 203(4), 1827–1839. <https://doi.org/10.1534/genetics.115.180992>

## Data Accessibility

*Structure\_threader*, the manual and example datasets are available on [https://github.com/StuntsPT/Structure\\_threader](https://github.com/StuntsPT/Structure_threader).

## Author Contributions

F. Pina-Martins has conceived the concept of the study, written most of the program code and written the manuscript. D. Silva has contributed to the ideas of the software, written code and revised the manuscript. J. Fino has contributed to the ideas of the software, written code and revised the manuscript. O. S. Paulo has mentored the work and revised the manuscript.

## Supporting Information

Spreadsheet with all run times (Supplementary Table 1).

7.

## Tables

*Table 1: Characteristics of the systems where the programs were benchmarked, along with the run time of the single threaded run.*

System Name	Type	CPU	Physical cores	Logical cores	OS	STRUCTURE single thread run time (s)	FASTSTRUCTURE single thread run time (s)	Maverick single thread run time (s)
		Frequency Base/Turbo (GHz)						
Haswell Laptop	i7 4700MQ	2.4/3.4	4	8	ArchLinux	9668	3140	1009
Ivy Bridge Desktop	i5 3350P	3.1/3.3	4	4	ArchLinux	10926	2854	1140
Nehalem Rack	Xeon E5520x2	2.26/2.53	8	16	Ubuntu 16.04	16000	6019	1835
Sandy Bridge Rack	Xeon E5-2609x2	2.4	8	8	Ubuntu 12.04	15805	5054	1711

# Figures




Figure 1: Threading strategy used in `Structure_threader`. Values in ellipses are read from the command line and passed to the main process, which generates a job queue. The jobs in the queue are then processed by the spawned child processes. The main process is responsible for handling and logging any errors that occur in the child processes.



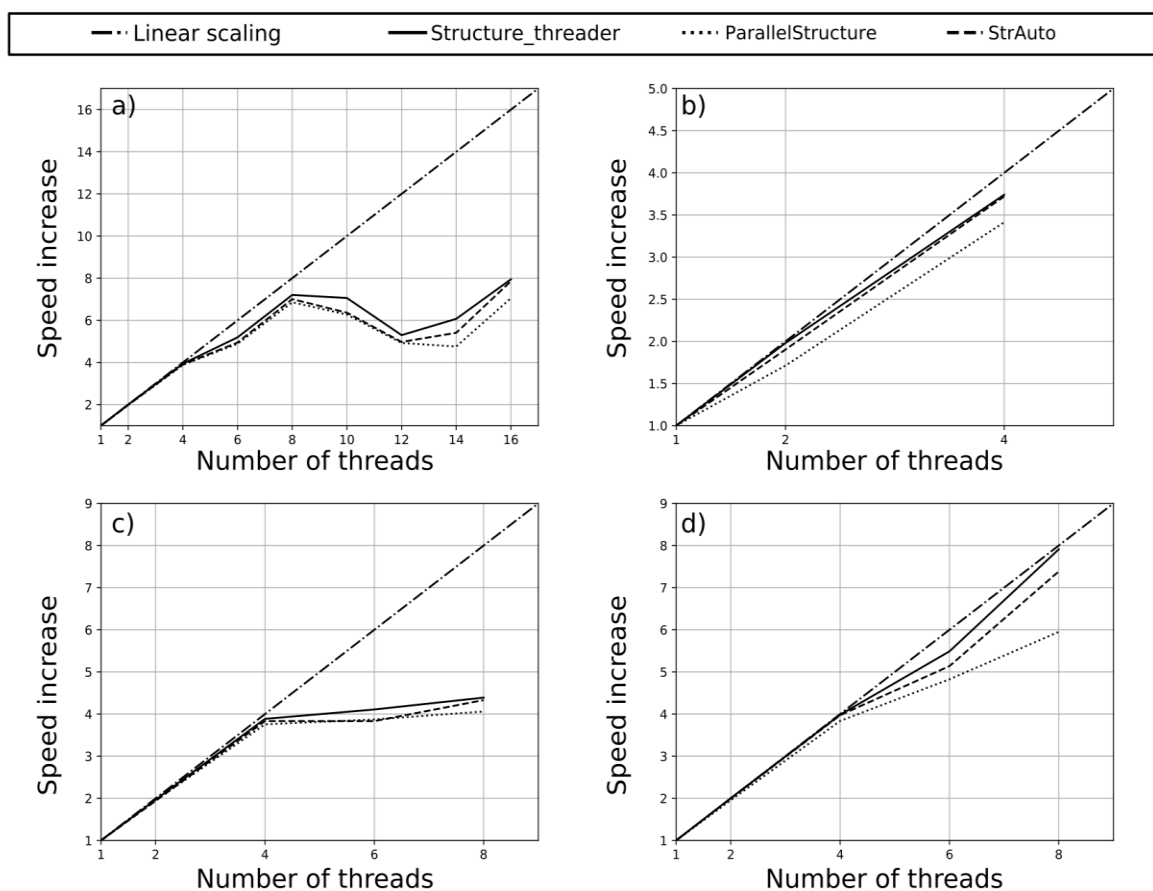


Figure 2: Plots of the STRUCTURE “speed up” as more threads are used in Structure\_threader, ParallelStructure and StrAuto. Each plot represents a different system – a) is “Nehalem Rack”, b) is “Ivy Bridge Desktop”, c) is “Haswell Laptop” and d) is “Sandy Bridge Rack”.

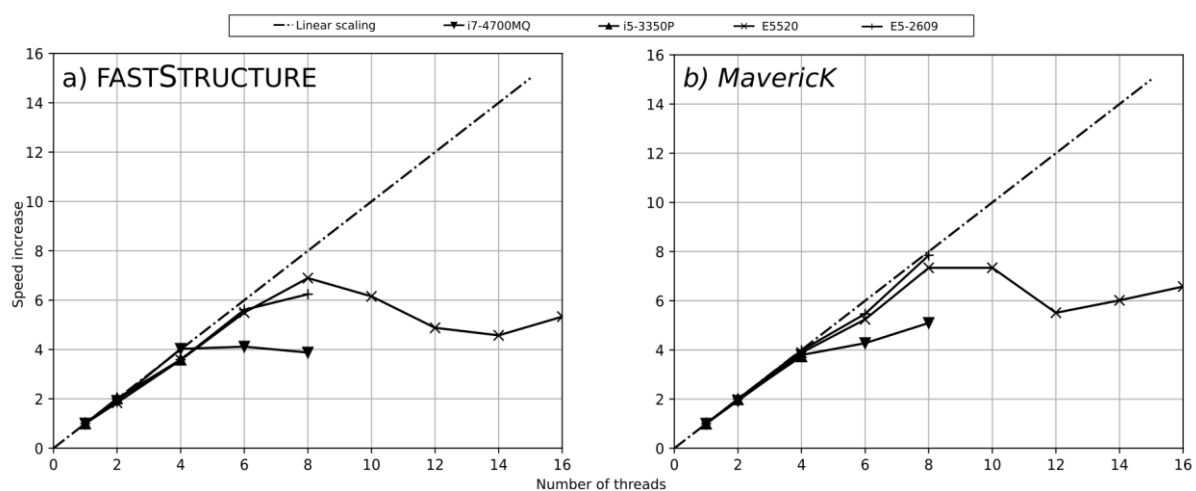


Figure 3: Plots of FASTSTRUCTURE and Maverick "speed up" as more threads are used when wrapped under Structure\_threader.